



INTRODUCTION TO MACHINE LEARNING

Clustering with k-means

Clustering, what?

- Cluster: collection of objects
 - *Similar* within cluster
 - *Dissimilar* between clusters
- Clustering: grouping objects in clusters
 - No labels: *unsupervised* classification
 - Plenty possible clusterings

Clustering, why?

- Pattern Analysis
- Visualise Data
- pre-Processing Step
- Outlier Detection
- ...
- Targeted Marketing Programs
- Student Segmentations
- Data Mining
- ...

Clustering, how?

- **Measure of Similarity: $d(..., ...)$**
 - Numerical variables → Metrics: Euclidean, Manhattan, ...
 - Categorical variables → Construct your own distance
- **Clustering Methods**
 - k-means
 - Hierarchical ← **Many variations**
 - ...

Compactness and Separation

- Within Cluster Sums of Squares (WSS):

$$WSS = \sum_{i=1}^{N_C} \sum_{x \in C_i} d(\mathbf{x}, \bar{\mathbf{x}}_{C_i})^2$$

Measure of compactness

$\bar{\mathbf{x}}_{C_i}$ Cluster Centroid
 \mathbf{x} Object
 C_i Cluster
 N_C #Clusters

Minimise WSS

- Between Cluster Sums of Squares (BSS):

$$BSS = \sum_{i=1}^{N_C} |C_i| \cdot d(\bar{\mathbf{x}}_{C_i}, \bar{\mathbf{x}})^2$$

Measure of separation


$\bar{\mathbf{x}}_{C_i}$ Cluster Centroid
 N_C #Clusters
 $|C_i|$ #Objects in Cluster
 $\bar{\mathbf{x}}$ Sample Mean

Maximise BSS

k-Means Algorithm

Goal: Partition data in k *disjoint* subsets

Let's take $k = 3$




k-Means Algorithm

Goal: Partition data in k *disjoint* subsets

1. Randomly assign k *centroids*

$k = 3$




k-Means Algorithm

Goal: Partition data in k *disjoint* subsets

1. Randomly assign k *centroids*
2. Assign data to *closest* centroid

$k = 3$




k-Means Algorithm

Goal: Partition data in k *disjoint* subsets

1. Randomly assign k *centroids*
2. Assign data to *closest* centroid
3. Moves centroids to *average* location

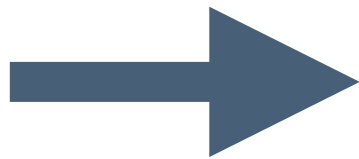
$k = 3$



Scale Issues

Metrics are often scale dependent!

Which pair is most similar? (Age, Income, IQ)

- $X_1 = (28, 72000, 120)$
 - $X_2 = (56, 73000, 80)$
 - $X_3 = (29, 74500, 118)$
- 
- Intuition: (X_1, X_3)
 - Euclidean: (X_1, X_2)

Solution: Rescale income / 1000\$

Standardizing

Problem: Multiple variables on different scales

Solution: Standardize your data

1. Subtract the mean
2. Divide by the standard deviation

```
> scale(data)
```

Note: Standardizing → Different interpretation



INTRODUCTION TO MACHINE LEARNING

Let's practice!



INTRODUCTION TO MACHINE LEARNING

Hierarchical Clustering

Hierarchical Clustering

Hierarchy:

- Which objects cluster first?
- Which cluster pairs merge? When?

Bottom-up:

- Starts from the objects
- Builds a hierarchy of clusters

Bottom-Up: Algorithm

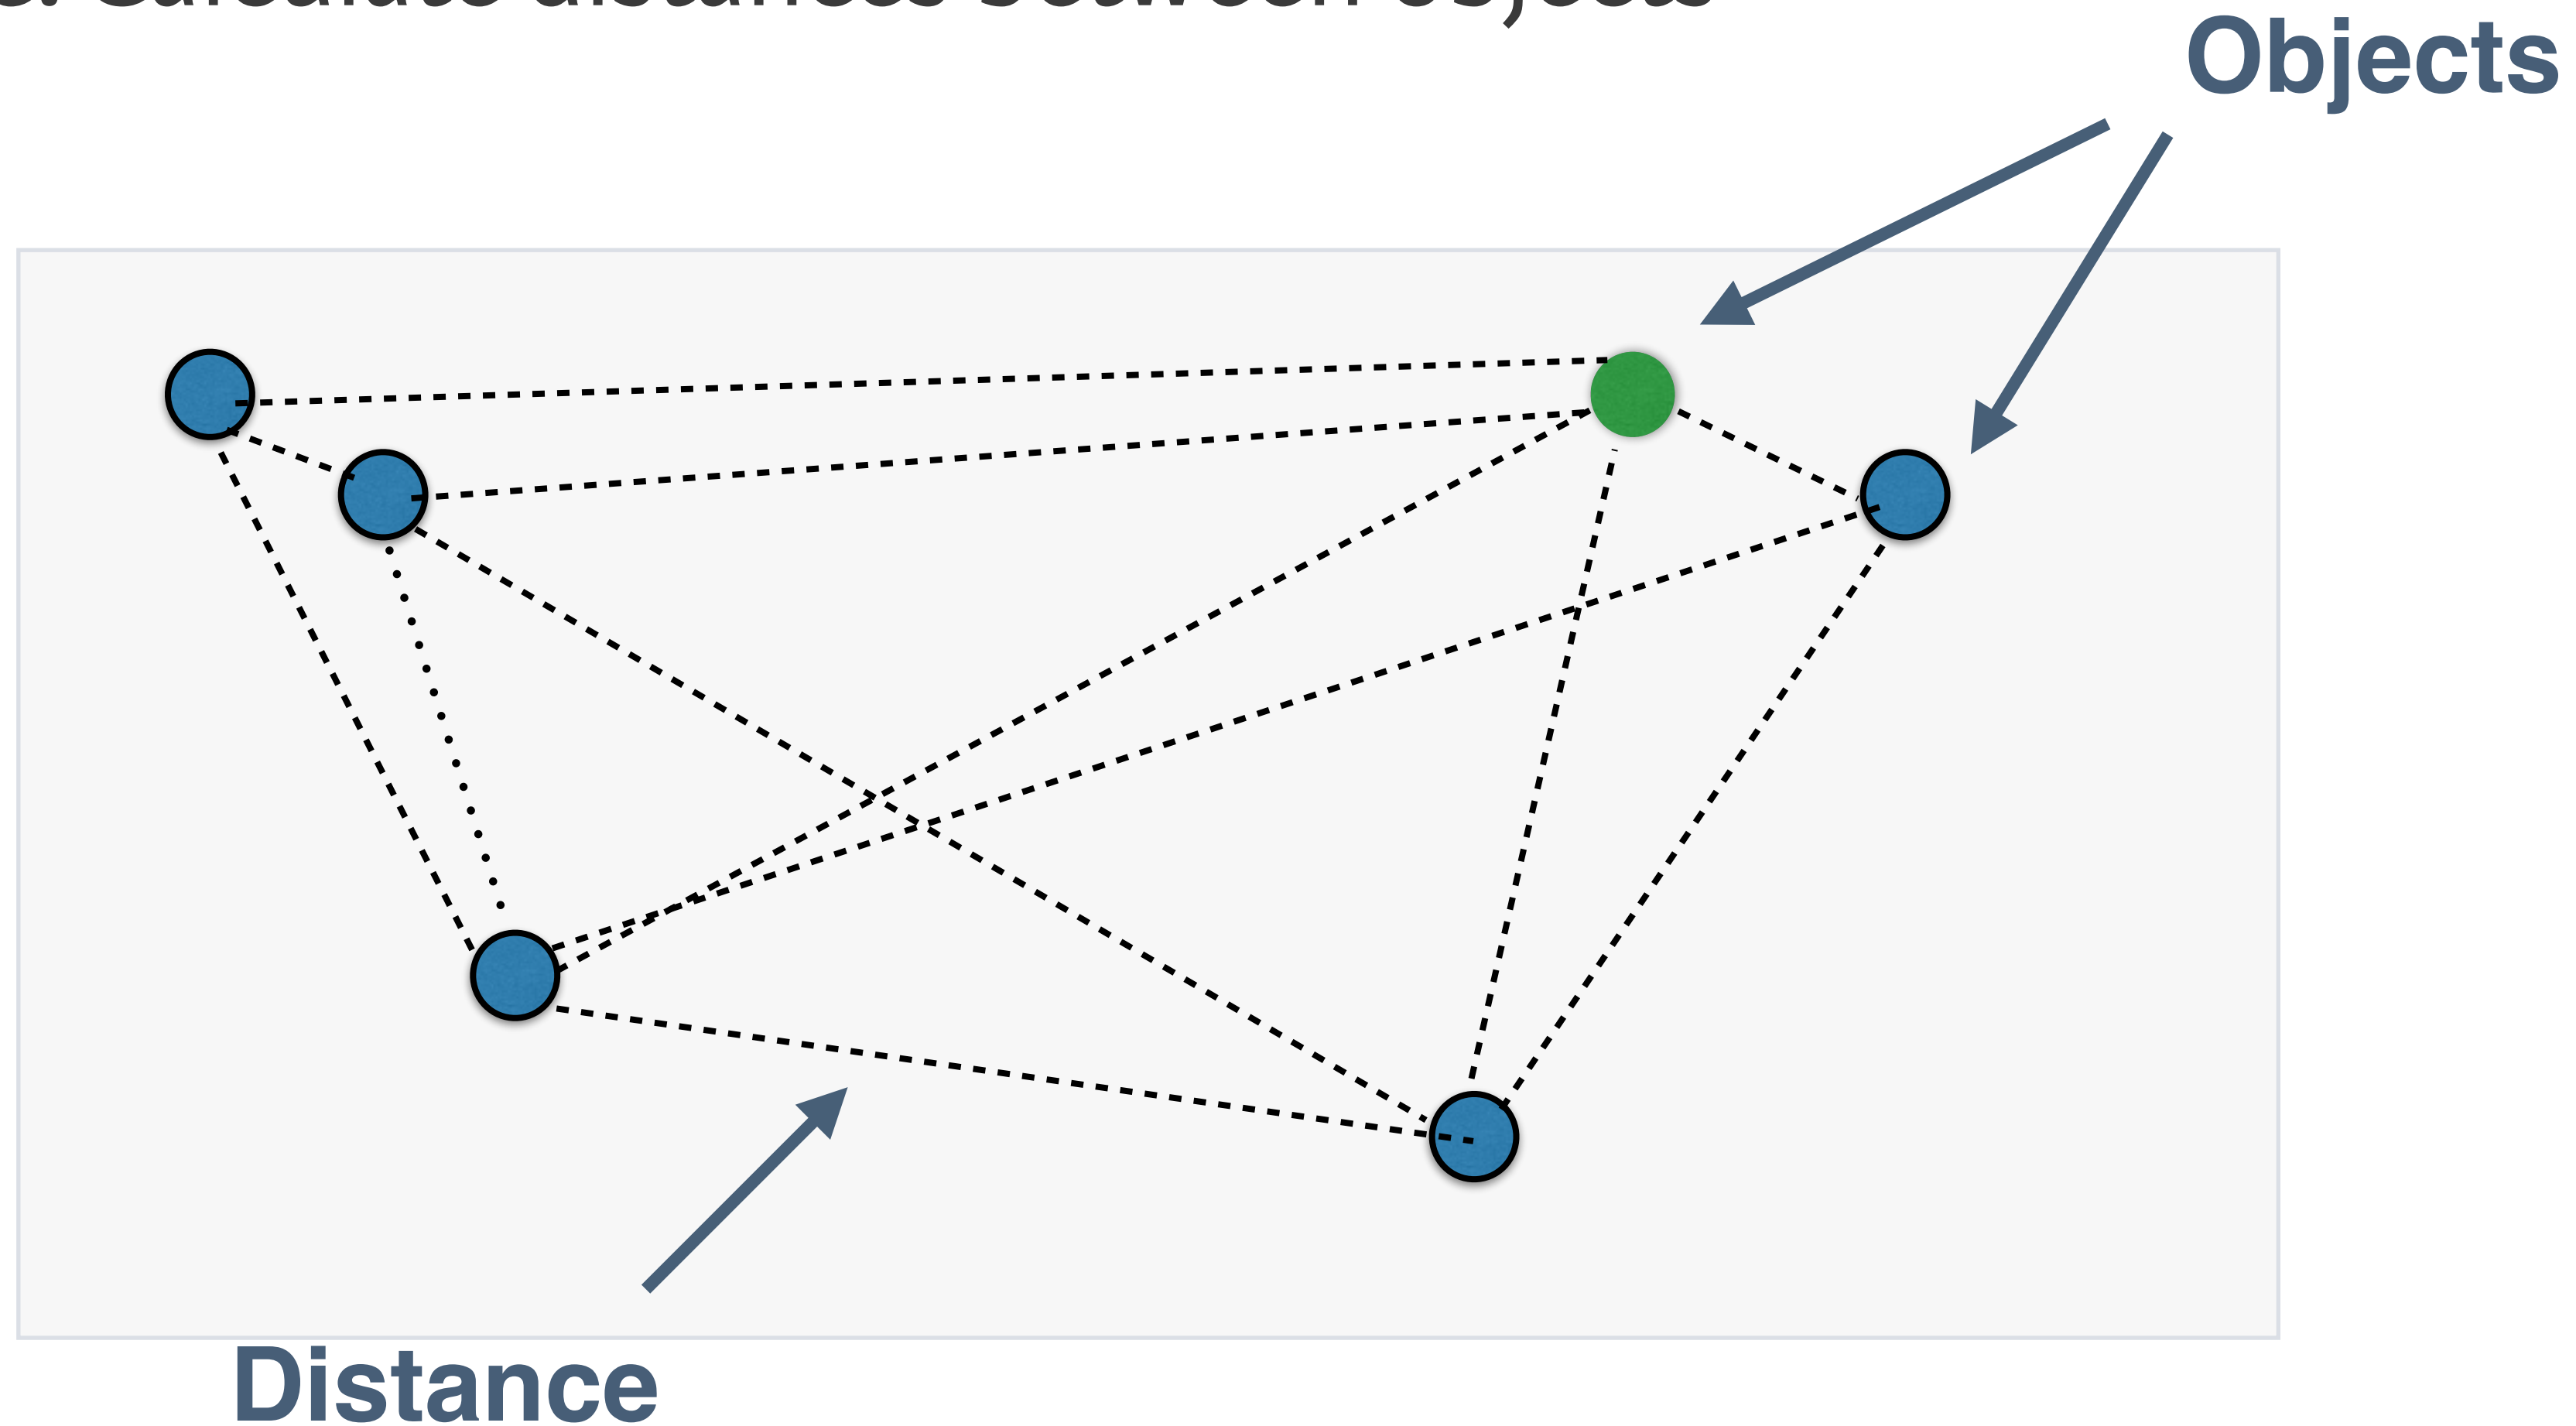
Pre: Calculate distances between objects

Objects



Bottom-Up: Algorithm

Pre: Calculate distances between objects



Bottom-Up: Algorithm

1. Put every object in its own cluster




Bottom-Up: Algorithm

2. Find the closest pair of clusters → Merge them



Bottom-Up: Algorithm

3. Compute distances between new cluster and old ones




Bottom-Up: Algorithm

4. Repeat steps two and three




Bottom-Up: Algorithm

4. Repeat steps two and three




Bottom-Up: Algorithm

4. Repeat steps two and three



Bottom-Up: Algorithm

4. Repeat steps two and three → One cluster




Linkage-Methods

- Simple-Linkage: *minimal* distance between clusters
- Complete-Linkage: *maximal* distance between clusters
- Average-Linkage: *average* distance between clusters

 Different Clusterings

Simple-Linkage

Minimal distance between objects in each clusters



k-Means: Pro and Cons

- Pros
 - Can undo merges
 - Fast computations
- Cons
 - Fixed #Clusters
 - Dependent on starting centroids



INTRODUCTION TO MACHINE LEARNING

Let's practice!