



MACHINE LEARNING TOOLBOX

Logistic regression on Sonar

Classification models

- Categorical (i.e. qualitative) target variable
- Example: will a loan default?
- Still a form of supervised learning
- Use a train/test split to evaluate performance
- Use the Sonar dataset
- Goal: distinguish rocks from mines

Example: Sonar data

```
> # Load the Sonar dataset  
> library(mlbench)  
> data(Sonar)
```

```
> # Look at the data
```

```
> Sonar[1:6, c(1:5, 61)]
```

	V1	V2	V3	V4	V5	Class
1	0.0200	0.0371	0.0428	0.0207	0.0954	R
2	0.0453	0.0523	0.0843	0.0689	0.1183	R
3	0.0262	0.0582	0.1099	0.1083	0.0974	R
4	0.0100	0.0171	0.0623	0.0205	0.0205	R
5	0.0762	0.0666	0.0481	0.0394	0.0590	R
6	0.0286	0.0453	0.0277	0.0174	0.0384	R

Splitting the data

- Randomly split data into training and test sets
- Use a 60/40 split, instead of 80/20
- Sonar dataset is small, so 60/40 gives a larger, more reliable test set

Splitting the data

```
# Randomly order the dataset
> rows <- sample(nrow(Sonar))
> Sonar <- Sonar[rows, ]

# Find row to split on
> split <- round(nrow(Sonar) * .60)
> train <- Sonar[1:split, ]
> test <- Sonar[(split + 1):nrow(Sonar), ]

# Confirm test set size
> nrow(train) / nrow(Sonar)
[1] 0.6009615
```



MACHINE LEARNING TOOLBOX

Let's practice!



MACHINE LEARNING TOOLBOX

Confusion matrix

Confusion matrix

Reference

Prediction

	Yes	No
Yes	True positive	False positive
No	False negative	True negative

Confusion matrix

```
# Fit a model
> model <- glm(Class ~ ., family = binomial(link = "logit"),
train)
> p <- predict(model, test, type = "response")
> summary(p)
   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
 0.0000  0.0000  0.9885  0.5296  1.0000  1.0000

# Turn probabilities into classes and look at their frequencies
> p_class <- ifelse(p > .50, "M", "R")
> table(p_class)
p_class
 M  R
44 39
```

Confusion matrix

- Make a 2-way frequency table
- Compare predicted vs. actual classes

```
# Make simple 2-way frequency table  
> table(p_class, test[["Class"]])  
p_class  M  R  
      M 13 31  
      R 30  9
```

Confusion matrix

```
# Use caret's helper function to calculate additional statistics  
> confusionMatrix(p_class, test[["Class"]])
```

```
      Reference  
Prediction  M   R  
      M  13  31  
      R  30   9
```

```
      Accuracy : 0.2651
```

```
      95% CI : (0.1742, 0.3734)
```

```
    No Information Rate : 0.5181
```

```
    P-Value [Acc > NIR] : 1
```

```
      Kappa : -0.4731
```

```
McNemar's Test P-Value : 1
```

```
      Sensitivity : 0.3023
```

```
      Specificity : 0.2250
```

```
    Pos Pred Value : 0.2955
```

```
    Neg Pred Value : 0.2308
```



MACHINE LEARNING TOOLBOX

Let's practice!



MACHINE LEARNING TOOLBOX

Class probabilities and class predictions

Different thresholds

- Not limited to 50% threshold
 - 10% would catch more mines with less certainty
 - 90% would catch fewer mines with more certainty
- Balance true positive and false positive rates
- Cost-benefit analysis

Confusion matrix

```
# Use a larger cutoff
> p_class <- ifelse(p > .99, "M", "R")
> table(p_class)
p_class
 M  R
41 42

# Make simple 2-way frequency table
> table(p_class, test[["Class"]])
p_class  M  R
      M 13 28
      R 30 12
```

Confusion matrix with caret

```
# Use caret to produce confusion matrix  
> confusionMatrix(p_class, test[["Class"]])
```

```
      Reference  
Prediction  M   R  
      M  13  28  
      R  30  12
```

```
      Accuracy : 0.3012
```

```
      95% CI : (0.2053, 0.4118)
```

```
      No Information Rate : 0.5181
```

```
      P-Value [Acc > NIR] : 1.0000
```

```
      Kappa : -0.397
```

```
      McNemar's Test P-Value : 0.8955
```

```
      Sensitivity : 0.3023
```

```
      Specificity : 0.3000
```

```
      Pos Pred Value : 0.3171
```

```
      Neg Pred Value : 0.2857
```




MACHINE LEARNING TOOLBOX

Let's practice!



MACHINE LEARNING TOOLBOX

Introducing the ROC curve

The challenge

- Many possible classification thresholds
- Requires manual work to choose
- Easy to overlook a particular threshold
- Need a more systematic approach


ROC curves

- Plot true/false positive rate at every possible threshold
- Visualize tradeoffs between two extremes **100% true positive rate vs. 0% false positive rate**
- Result is an ROC curve
- Developed as a method for analyzing radar signals

An example ROC curve

```
# Create ROC curve  
> library(caTools)  
> colAUC(p, test[["Class"]], plotROC = TRUE)
```

- X-axis: false positive rate
- Y-axis: true positive rate
- Each point along the curve represents a different threshold





MACHINE LEARNING TOOLBOX


Let's practice!



MACHINE LEARNING TOOLBOX

Area under the curve (AUC)

From ROC to AUC



Defining AUC

- Single-number summary of model accuracy
- Summarizes performance across all thresholds
- Rank different models within the same dataset

Defining AUC

- Ranges from 0 to 1
 - 0.5 = random guessing
 - 1 = model always right
 - 0 = model always wrong
- Rule of thumb: AUC as a letter grade
 - 0.9 = "A"
 - 0.8 = "B"
 - ...



MACHINE LEARNING TOOLBOX

Let's practice!