



MACHINE LEARNING TOOLBOX

Median imputation

Dealing with missing values

- Most models require numbers, can't handle missing data
- Common approach: remove rows with missing data
 - Can lead to biases in data
 - Generate over-confident models
- Better strategy: median imputation!
 - Replace missing values with medians
 - Works well if data *missing at random* (MAR)

Example: mtcars

```
# Generate some data with missing values
> data(mtcars)
> set.seed(42)
> mtcars[sample(1:nrow(mtcars), 10), "hp"] <- NA

# Split target from predictors
> Y <- mtcars$mpg
> X <- mtcars[, 2:4]

# Try to fit a caret model
> library(caret)
> model <- train(x = X, y = Y)
Error in train.default(x = X, y = Y) : Stopping
```

A simple solution

```
# Now fit with median imputation  
> model <- train(x = X, y = Y, preProcess = "medianImpute")  
> print(model)  
Random Forest
```

```
32 samples  
 3 predictor
```

```
Pre-processing: median imputation (3)  
Resampling: Bootstrapped (25 reps)  
Summary of sample sizes: 32, 32, 32, 32, 32, 32, ...  
Resampling results across tuning parameters:
```

| mtry | RMSE | Rsquared |
|------|----------|-----------|
| 2 | 2.617096 | 0.8234652 |
| 3 | 2.670550 | 0.8164535 |

RMSE was used to select the optimal model using the smallest value.
The final value used for the model was mtry = 2.



MACHINE LEARNING TOOLBOX

Let's practice!



MACHINE LEARNING TOOLBOX

KNN imputation

Dealing with missing values

- Median imputation is fast, but...
- Can produce incorrect results if data *missing not at random*
- k-nearest neighbors (KNN) imputation
- Imputes based on "similar" non-missing rows

Example: missing not at random

- Pretend smaller cars don't report horsepower
- Median imputation incorrect in this case **Assumes small cars have medium-large horsepower**

```
# Generate data with missing values
> data(mtcars)
> mtcars[mtcars$disp < 140, "hp"] <- NA
> Y <- mtcars$mpg
> X <- mtcars[, 2:4]

# Use median imputation
> set.seed(42)
> model <- train(x = X, y = Y, method = "glm",
                 preProcess = "medianImpute")
> print(min(model$results$RMSE))
[1] 3.612713
```


Example: missing not at random

- KNN imputation is better
- Uses cars with similar `disp` / `cyl` to impute
- Yields a more accurate (but slower) model

```
# Use KNN imputation
> set.seed(42)
> model <- train(x = X, y = Y,
                 method = "glm",
                 preProcess = "knnImpute"
               )
> print(min(model$results$RMSE))
[1] 3.558881 Compare to 3.61 for median imputation
```



MACHINE LEARNING TOOLBOX

Let's practice!



MACHINE LEARNING TOOLBOX

Multiple preprocessing methods

The wide world of preProcess

- You can do a lot more than median or knn imputation!
- Can chain together multiple preprocessing steps
- Common "recipe" for linear models (order matters!)
Median imputation -> center -> scale -> fit glm
- See ?preProcess for more detail

Example: preprocessing mtcars

```
# Generate some data with missing values
> data(mtcars)
> set.seed(42)
> mtcars[sample(1:nrow(mtcars), 10), "hp"] <- NA
> Y <- mtcars$mpg
> X <- mtcars[,2:4]      Missing at random

# Use linear model "recipe"
> set.seed(42)
> model <- train(
  x = X, y = Y, method = "glm",
  preProcess = c("medianImpute", "center", "scale")
)
> print(min(model$results$RMSE))
[1] 3.612713
```

Example: preprocessing mtcars

```
# PCA before modeling
> set.seed(42)
> model <- train(
  x = X, y = Y, method = "glm",
  preProcess = c("medianImpute", "center", "scale", "pca")
)
> min(model$results$RMSE)
[1] 3.402557
```

Example: preprocessing mtcars

```
# Spatial sign transform
> set.seed(42)
> model <- train(
  x = X, y = Y, method = "glm",
  preProcess = c("medianImpute", "center", "scale", "spatialSign"))
> min(model$results$RMSE)
[1] 4.284904
```

Preprocessing cheat sheet

- Start with median imputation **Try KNN imputation if data missing not at random**
- For linear models...
 - Center and scale
 - Try PCA and spatial sign
- Tree-based models don't need much preprocessing



MACHINE LEARNING TOOLBOX

Let's practice!



MACHINE LEARNING TOOLBOX

Handling low-information predictors

No (or low) variance variables

- Some variables don't contain much information
 - Constant (i.e. no variance)
 - Nearly constant (i.e. low variance)
- Easy for one fold of CV to end up with constant column
- Can cause problems for your models
- Usually remove extremely low variance variables

Example: constant column in mtcars

```
# Reproduce dataset from last video
> data(mtcars)
> set.seed(42)
> mtcars[sample(1:nrow(mtcars), 10), "hp"] <- NA
> Y <- mtcars$mpg
> X <- mtcars[, 2:4]

# Add constant-valued column to mtcars
> X$bad <- 1
```

Example: constant column in mtcars

```
# Try to fit a model with PCA + glm
> model <- train(
  x = X, y = Y, method = "glm",
  preProcess = c("medianImpute", "center", "scale", "pca")
)
```

Warning in preProcess.default(thresh = 0.95, k = 5, method =
c("medianImpute", :

These variables have zero variances: bad

Something is wrong; all the RMSE metric values are missing:

| RMSE | | Rsquared | |
|----------|------|----------|------|
| Min. | : NA | Min. | : NA |
| 1st Qu.: | NA | 1st Qu.: | NA |
| Median | : NA | Median | : NA |
| Mean | :NaN | Mean | :NaN |
| 3rd Qu.: | NA | 3rd Qu.: | NA |
| Max. | : NA | Max. | : NA |
| NA's | :1 | NA's | :1 |

caret to the rescue (again)

- "zv" removes constant columns
- "nzv" removes nearly constant columns

```
# Have caret remove those columns during modeling
> set.seed(42)
> model <- train(
  x = X, y = Y, method = "glm",
  preProcess = c("zv", "medianImpute", "center", "scale", "pca")
)
> min(model$results$RMSE)
[1] 3.402557
```



MACHINE LEARNING TOOLBOX

Let's practice!



MACHINE LEARNING TOOLBOX


Principle components analysis (PCA)

Principle components analysis

- Combines low-variance and correlated variables
- Single set of high-variance, perpendicular predictors
- Prevents collinearity (i.e. correlation among predictors)

PCA: a visual representation

- First component has highest variance
- Second component has second highest variance
- And so on...



Example: blood-brain data

- Lots of predictors
- Many of them low-variance

```
# Load the blood brain dataset
> data(BloodBrain)
> names(bbbDescr)[nearZeroVar(bbbDescr)]
[1] "negative"      "peoe_vsa.2.1" "peoe_vsa.3.1" "a_acid"
[5] "vsa_acid"      "frac.anion7." "alert"
```

Example: blood-brain data

```
# Basic model
> set.seed(42)
> data(BloodBrain)
> model <- train(
  x = bbbDescr, y = logBBB, method = "glm",
  trControl = trainControl(method = "cv", number = 10, verbose = TRUE),
  preProcess = c("zv", "center", "scale")
)
> min(model$results$RMSE)
[1] 1.107702
```

Example: blood-brain data

```
# Remove low-variance predictors
> set.seed(42)
> data(BloodBrain)
> model <- train(
  x = bbbDescr, y = logBBB, method = "glm",
  trControl = trainControl(method = "cv", number = 10, verbose = TRUE),
  preProcess = c("nzv", "center", "scale")
)
> min(model$results$RMSE)
[1] 0.9796199
```

Example: blood-brain data

```
# Add PCA
> set.seed(42)
> data(BloodBrain)
> model <- train(
  x = bbbDescr, y = logBBB, method = "glm",
  trControl = trainControl(method = "cv", number = 10, verbose = TRUE),
  preProcess = c("zv", "center", "scale", "pca")
)
> min(model$results$RMSE)
[1] 0.9796199
```



MACHINE LEARNING TOOLBOX

Let's practice!