



MACHINE LEARNING WITH THE EXPERTS: SCHOOL BUDGETS

**It's
time to build a
model**



It's time to build a model

- Always a good approach to start with a very simple model
- Gives a sense of how challenging the problem is
- Many more things can go wrong in complex models
- How much signal can we pull out using basic methods?

It's time to build a model

- Train basic model on numeric data only
 - Want to go from raw data to predictions quickly
- Multi-class logistic regression
 - Train classifier on each label separately and use those to predict
- Format predictions and save to csv
- Compute log loss score

Splitting the multi-class dataset

- Recall: Train-test split
 - Will not work here
 - May end up with labels in test set that never appear in training set
- Solution: StratifiedShuffleSplit
 - Only works with a single target variable
 - We have many target variables
 - `multilabel_train_test_split()`



Training the model

```
In [4]: from sklearn.linear_model import LogisticRegression  
  
In [5]: from sklearn.multiclass import OneVsRestClassifier  
  
In [6]: clf = OneVsRestClassifier(LogisticRegression())  
  
In [7]: clf.fit(X_train, y_train)
```

- OneVsRestClassifier:
 - Treats each column of y independently
 - Fits a separate classifier for each of the columns



MACHINE LEARNING WITH THE EXPERTS: SCHOOL BUDGETS

Let's practice!



MACHINE LEARNING WITH THE EXPERTS: SCHOOL BUDGETS

Making predictions



Predicting on holdout data

```
In [1]: holdout = pd.read_csv('HoldoutData.csv', index_col=0)
```

```
In [2]: holdout = holdout[NUMERIC_COLUMNS].fillna(-1000)
```

```
In [3]: predictions = clf.predict_proba(holdout)
```

- If `.predict()` was used instead:
 - Output would be 0 or 1
 - Log loss penalizes being confident and wrong
 - Worse performance compared to `.predict_proba()`



Submitting your predictions as a csv

| | Function <input type="checkbox"/> Aides Compensation | Function <input type="checkbox"/> Career & Academic Counseling | Function <input type="checkbox"/> Communications | ... | Use <input type="checkbox"/> O&M | Use <input type="checkbox"/> Pupil Services & Enrichment | Use <input type="checkbox"/> Untracked Budget Set- Aside |
|--------|---|--|--|-----|----------------------------------|--|--|
| 180042 | 0.027027 | 0.027027 | 0.027027 | ... | 0.125 | 0.125 | 0.125 |
| 28872 | 0.027027 | 0.027027 | 0.027027 | ... | 0.125 | 0.125 | 0.125 |
| 186915 | 0.027027 | 0.027027 | 0.027027 | ... | 0.125 | 0.125 | 0.125 |
| 412396 | 0.027027 | 0.027027 | 0.027027 | ... | 0.125 | 0.125 | 0.125 |
| 427740 | 0.027027 | 0.027027 | 0.027027 | ... | 0.125 | 0.125 | 0.125 |

- All formatting can be done with the pandas `to_csv` function

Format and submit predictions

```
In [4]: prediction_df = pd.DataFrame(columns=pd.get_dummies(df[LABELS],  
...:                                     prefix_sep='__').columns,  
...:                                     index=holdout.index,  
...:                                     data=predictions)
```

```
In [5]: prediction_df.to_csv('predictions.csv')
```

```
In [6]: score = score_submission(pred_path='predictions.csv')
```



DrivenData leaderboard

| | User or team | Public ⓘ ⬆ | Private ⬆ | Timestamp ⓘ | Trend ⓘ ⬆ | # Entries ⬆ |
|--|--------------------|------------|-----------|---------------------------|-----------|-------------|
| | quocnle | 0.3665 | 0.3650 | Jan. 6, 2015, 12:27 a.m. | | 96 |
| | Abhishek | 0.4409 | 0.4388 | Jan. 6, 2015, 4:09 p.m. | | 71 |
| | giba | 0.4551 | 0.4534 | Jan. 5, 2015, 4:52 p.m. | | 34 |
| | trev | 0.5054 | 0.5001 | Jan. 3, 2015, 2 a.m. | | 23 |
| | Kappa | 0.5228 | 0.5195 | Jan. 6, 2015, 11:46 p.m. | | 17 |
| | bamine | 0.5344 | 0.5298 | Dec. 12, 2014, 12:52 a.m. | | 39 |
| | futuristic reality | 0.5512 | 0.5477 | Nov. 24, 2014, 8:54 a.m. | | 22 |
| | JesseBuesking | 0.5584 | 0.5556 | Jan. 6, 2015, 4:51 p.m. | | 15 |
| | mkrump | 0.5817 | 0.5769 | Jan. 3, 2015, 5:12 p.m. | | 57 |
| | joel314 | 0.5806 | 0.5772 | Dec. 10, 2014, 4:41 p.m. | | 63 |



MACHINE LEARNING WITH THE EXPERTS: SCHOOL BUDGETS

Let's practice!



MACHINE LEARNING WITH THE EXPERTS: SCHOOL BUDGETS

A very brief introduction to NLP



A very brief introduction to NLP

- Data for NLP:
 - Text, documents, speech, ...
- Tokenization
 - Splitting a string into segments
 - Store segments as list
- Example: 'Natural Language Processing'
 - —> ['Natural', 'Language', 'Processing']



Tokens and token patterns

- Tokenize on whitespace

PETRO-VEND FUEL AND FLUIDS

PETRO-VEND | FUEL | AND | FLUIDS

- Tokenize on whitespace *and* punctuation

PETRO-VEND FUEL AND FLUIDS


PETRO | VEND | FUEL | AND | FLUIDS

Bag of words representation

- Count the number of times a particular token appears
- “Bag of words”
 - Count the number of times a word was pulled out of the bag
- This approach discards information about word order
 - “Red, not blue” is the same as “blue, not red”



1-gram, 2-gram, ..., n-gram





MACHINE LEARNING WITH THE EXPERTS: SCHOOL BUDGETS

Let's practice!



MACHINE LEARNING WITH THE EXPERTS: SCHOOL BUDGETS

Representing text numerically



Representing text numerically

- Bag-of-words
 - Simple way to represent text in machine learning
 - Discards information about grammar and word order
 - Computes frequency of occurrence



Scikit-learn tools for bag-of-words

- `CountVectorizer()`
 - Tokenizes all the strings
 - Builds a ‘vocabulary’
 - Counts the occurrences of each token in the vocabulary



Using CountVectorizer() on column of main dataset

```
In [1]: from sklearn.feature_extraction.text import CountVectorizer
```

```
In [2]: TOKENS_BASIC = '\\S+(?=?\\s+)'
```

```
In [3]: df.Program_Description.fillna('', inplace=True)
```

```
In [4]: vec_basic = CountVectorizer(token_pattern=TOKENS_BASIC)
```



Using CountVectorizer() on column of main dataset

```
In [5]: vec_basic.fit(df.Program_Description)
```

```
Out[5]:
```

```
CountVectorizer(analyzer='word', binary=False, decode_error='strict',  
                dtype=<class 'numpy.int64'>, encoding='utf-8', input='content',  
                lowercase=True, max_df=1.0, max_features=None, min_df=1,  
                ngram_range=(1, 1), preprocessor=None, stop_words=None,  
                strip_accents=None, token_pattern='\\S+(?=\\s+)',  
                tokenizer=None, vocabulary=None)
```

```
In [6]: msg = 'There are {} tokens in Program_Description if tokens are  
any non-whitespace'
```

```
In [7]: print(msg.format(len(vec_basic.get_feature_names())))
```

```
There are 157 tokens in Program_Description if tokens are any non-  
whitespace
```




MACHINE LEARNING WITH THE EXPERTS: SCHOOL BUDGETS

Let's practice!