



MACHINE LEARNING WITH THE EXPERTS: SCHOOL BUDGETS

Pipelines, feature & text preprocessing



The pipeline workflow

- Repeatable way to go from raw data to trained model
- Pipeline object takes sequential list of steps
 - Output of one step is input to next step
- Each step is a tuple with two elements
 - Name: string
 - Transform: obj implementing `.fit()` and `.transform()`
- Flexible: a step can itself be another pipeline!

Instantiate simple pipeline with one step

```
In [1]: from sklearn.pipeline import Pipeline

In [2]: from sklearn.linear_model import LogisticRegression

In [3]: from sklearn.multiclass import OneVsRestClassifier

In [4]: pl = Pipeline([
...:     ('clf', OneVsRestClassifier(LogisticRegression()))
...:     ])
```

Train and test with sample numeric data

```
In [5]: sample_df.head()
```

```
Out[5]:
```

	label	numeric	text	with_missing
0	a	-4.167578	bar	-4.084883
1	a	-0.562668		2.043464
2	a	-21.361961		-33.315334
3	b	16.402708	foo bar	30.884604
4	a	-17.934356	foo	-27.488405



Train and test with sample numeric data

```
In [6]: from sklearn.model_selection import train_test_split
```

```
In [7]: X_train, X_test, y_train, y_test = train_test_split(  
...:                                     sample_df[['numeric']],  
...:                                     pd.get_dummies(sample_df['label']),  
...:                                     random_state=2)
```

```
In [8]: pl.fit(X_train, y_train)
```

```
Out[8]:
```

```
Pipeline(steps=[('clf', OneVsRestClassifier(estimator=LogisticRegression(C=1.0,  
class_weight=None, dual=False, fit_intercept=True,  
intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,  
penalty='l2', random_state=None, solver='liblinear', tol=0.0001,  
verbose=0, warm_start=False),  
n_jobs=1))])
```



Train and test with sample numeric data

```
In [9]: accuracy = pl.score(X_test, y_test)
```

```
In [10]: print('accuracy on numeric data, no nans: ', accuracy)
accuracy on numeric data, no nans: 0.44
```



Adding more steps to the pipeline

```
In [11]: X_train, X_test, y_train, y_test = train_test_split(sample_df[['numeric',  
...:                                                         'with_missing']], pd.get_dummies(  
...:                                                         sample_df['label']), random_state=2)
```

```
In [12]: pl.fit(X_train, y_train)
```

```
Traceback (most recent call last):
```

```
...
```

```
ValueError: Input contains NaN, infinity or a value too large for  
dtype('float64').
```



Preprocessing numeric features with missing data

```
In [13]: from sklearn.preprocessing import Imputer

In [14]: X_train, X_test, y_train, y_test = train_test_split(sample_df[['numeric',
...:                                                         'with_missing']], pd.get_dummies(
...:                                                         sample_df['label']), random_state=2)

In [15]: pl = Pipeline([
...:     ('imp', Imputer()),
...:     ('clf', OneVsRestClassifier(LogisticRegression()))
...:     ])
```




Preprocessing numeric features with missing data

```
In [16]: pipeline.fit(X_train, y_train)
```

```
In [17]: accuracy = pl.score(X_test, y_test)
```

```
In [18]: print('accuracy on all numeric, incl nans: ', accuracy)
accuracy on all numeric, incl nans: 0.48
```

- No errors!



MACHINE LEARNING WITH THE EXPERTS: SCHOOL BUDGETS

Let's practice!



MACHINE LEARNING WITH THE EXPERTS: SCHOOL BUDGETS

Text features and feature unions



Preprocessing text features

```
In [1]: from sklearn.feature_extraction.text import CountVectorizer
```

```
In [2]: X_train, X_test, y_train, y_test = train_test_split(sample_df['text'],  
...:                                                         pd.get_dummies(  
...:                                                         sample_df['label']),  
...:                                                         random_state=2)
```

```
In [3]: pl = Pipeline([  
...:     ('vec', CountVectorizer()),  
...:     ('clf', OneVsRestClassifier(LogisticRegression()))  
...: ])
```



Preprocessing text features

```
In [4]: pl.fit(X_train, y_train)
```

```
Out[4]:
```

```
Pipeline(steps=[('vec', CountVectorizer(analyzer='word', binary=False,
decode_error='strict', dtype=<class 'numpy.int64'>, encoding='utf-8',
input='content', lowercase=True, max_df=1.0, max_features=None, min_df=1,
ngram_range=(1, 1), preprocessor=None, stop_words=None, strip_...=None,
solver='liblinear', tol=0.0001, verbose=0, warm_start=False), n_jobs=1))])
```

```
In [5]: accuracy = pl.score(X_test, y_test)
```

```
In [6]: print('accuracy on sample data: ', accuracy)
```

```
accuracy on sample data: 0.64
```



Preprocessing multiple dtypes

- Want to use all available features in one pipeline
- Problem
 - Pipeline steps for numeric and text preprocessing can't follow each other
 - e.g., output of CountVectorizer can't be input to Imputer
- Solution
 - FunctionTransformer() & FeatureUnion()



FunctionTransformer

- Turns a Python function into an object that a scikit-learn pipeline can understand
- Need to write two functions for pipeline preprocessing
 - Take entire DataFrame, return numeric columns
 - Take entire DataFrame, return text columns
- Can then preprocess numeric and text data in separate pipelines



Putting it all together

```
In [7]: X_train, X_test, y_train, y_test = train_test_split(sample_df[['numeric',  
...:                                                         'with_missing', 'text']], pd.get_dummies(  
...:                                                         sample_df['label']), random_state=2)
```

```
In [8]: from sklearn.preprocessing import FunctionTransformer
```

```
In [9]: from sklearn.pipeline import FeatureUnion
```





FeatureUnion Text and Numeric Features


```
In [12]: from sklearn.pipeline import FeatureUnion
```

```
In [13]: union = FeatureUnion([  
...:     ('numeric', numeric_pipeline),  
...:     ('text', text_pipeline)  
...: ])
```

Text Features



Numeric Features





Putting it all together

```
In [14]: numeric_pipeline = Pipeline([
...:     ('selector', get_numeric_data),
...:     ('imputer', Imputer())
...: ])

In [15]: text_pipeline = Pipeline([
...:     ('selector', get_text_data),
...:     ('vectorizer', CountVectorizer())
...: ])

In [16]: pl = Pipeline([
...:     ('union', FeatureUnion([
...:         ('numeric', numeric_pipeline),
...:         ('text', text_pipeline)
...:     ])),
...:     ('clf', OneVsRestClassifier(LogisticRegression()))
...: ])
```



MACHINE LEARNING WITH THE EXPERTS: SCHOOL BUDGETS

Let's practice!



MACHINE LEARNING WITH THE EXPERTS

Choosing a classification model



Main dataset: lots of text

```
In [1]: LABELS = ['Function', 'Use', 'Sharing', 'Reporting', 'Student_Type',  
...:              'Position_Type', 'Object_Type', 'Pre_K', 'Operating_Status']
```

```
In [2]: NON_LABELS = [c for c in df.columns if c not in LABELS]
```

```
In [3]: len(NON_LABELS) - len(NUMERIC_COLUMNS)
```

```
Out[3]: 14
```



Using pipeline with the main dataset

```
In [4]: import numpy as np

In [5]: import pandas as pd

In [6]: df = pd.read_csv('TrainingSetSample.csv', index_col=0)

In [7]: dummy_labels = pd.get_dummies(df[LABELS])

In [8]: X_train, X_test, y_train, y_test = multilabel_train_test_split(
...:                                     df[NON_LABELS], dummy_labels,
...:                                     0.2)
```



Using pipeline with the main dataset

```
In [10]: get_text_data = FunctionTransformer(
...:                                     combine_text_columns,
...:                                     validate=False)

In [11]: get_numeric_data = FunctionTransformer(lambda x:
...:                                             x[NUMERIC_COLUMNS], validate=False)

In [12]: pl = Pipeline([
...:     ('union', FeatureUnion([
...:         ('numeric_features', Pipeline([
...:             ('selector', get_numeric_data),
...:             ('imputer', Imputer())
...:         ])),
...:         ('text_features', Pipeline([
...:             ('selector', get_text_data),
...:             ('vectorizer', CountVectorizer())
...:         ]))
...:     ])
...:     ),
...:     ('clf', OneVsRestClassifier(LogisticRegression()))
...: ])
```




Performance using main dataset

```
In [13]: pl.fit(X_train, y_train)
Out[13]:
Pipeline(steps=[('union', FeatureUnion(n_jobs=1,
transformer_list=[('numeric_features', Pipeline(steps=[('selector',
FunctionTransformer(accept_sparse=False, func=<function <lambda> at
0x11415ec80>, pass_y=False, validate=False)), ('imputer', Imputer(axis=0,
copy=True, missing_valu...=None, solver='liblinear', tol=0.0001, verbose=0,
warm_start=False),n_jobs=1)))]])
```

Flexibility of model step

- Is current model the best?
- Can quickly try different models with pipelines
 - Pipeline preprocessing steps unchanged
 - Edit the model step in your pipeline
 - Random Forest, Naïve Bayes, k-NN

Easily try new models using pipeline

```
In [14]: from sklearn.ensemble import RandomForestClassifier
```

```
In [15]: pl = Pipeline([
...:     ('union', FeatureUnion(
...:         transformer_list = [
...:             ('numeric_features', Pipeline([
...:                 ('selector', get_numeric_data),
...:                 ('imputer', Imputer())
...:             ])),
...:             ('text_features', Pipeline([
...:                 ('selector', get_text_data),
...:                 ('vectorizer', CountVectorizer())
...:             ]))
...:         ])
...:     ])
...:     ('clf', OneVsRest(RandomForestClassifier()))
...: ])
```



MACHINE LEARNING WITH THE EXPERTS: SCHOOL BUDGETS

Let's practice!