



UNSUPERVISED LEARNING IN PYTHON

Non-negative matrix factorization (NMF)



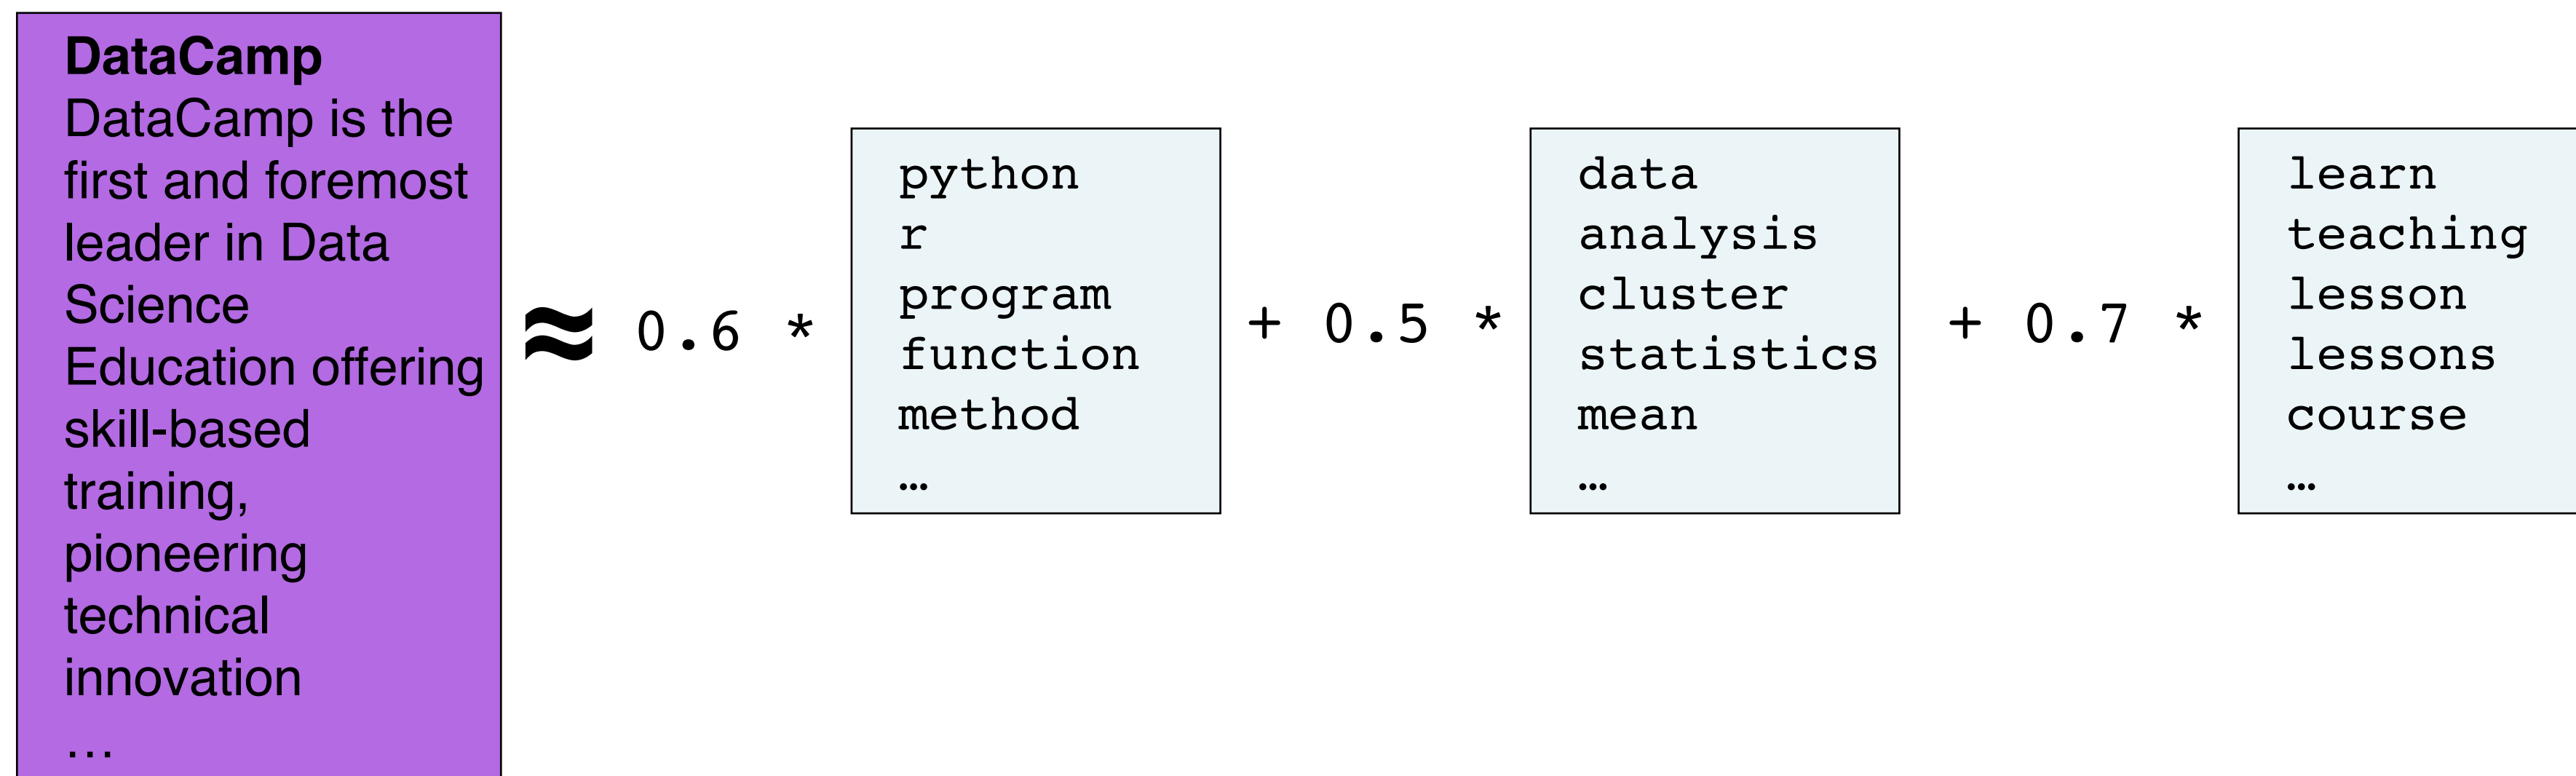
Non-negative matrix factorization

- NMF = "non-negative matrix factorization"
- Dimension reduction technique
- NMF models are *interpretable* (unlike PCA)
- Easy to interpret means easy to explain!
- However, all sample features must be non-negative (≥ 0)



Interpretable parts

- NMF expresses documents as combinations of topics (or "themes")





Interpretable parts

- NMF expresses images as combinations of patterns

$$\text{Image of '7'} \approx 0.98 * \text{Pattern 1} + 0.91 * \text{Pattern 2} + 0.94 * \text{Pattern 3}$$



Using `scikit-learn` NMF

- Follows `fit()` / `transform()` pattern
- Must specify number of components e.g. `NMF(n_components=2)`
- Works with NumPy arrays and with `csr_matrix`



Example word-frequency array

- Word frequency array, 4 words, many documents
- Measure presence of words in each document using "tf-idf"
- "tf" = frequency of word in document
- "idf" reduces influence of frequent words

	course	datacamp	potato	the
document0	0.2,	0.3,	0.0,	0.1
document1	0.0,	0.0,	0.4,	0.1
...			...	



Example usage of NMF

- `samples` is the word-frequency array

```
In [1]: from sklearn.decomposition import NMF
```

```
In [2]: model = NMF(n_components=2)
```

```
In [3]: model.fit(samples)
```

```
Out[3]: NMF(alpha=0.0, ... )
```

```
In [4]: nmf_features = model.transform(samples)
```



NMF components

- NMF has components
- ... just like PCA has principal components
- Dimension of components = dimension of samples
- Entries are non-negative

```
In [5]: print(model.components_)  
[[ 0.01  0.    2.13  0.54]  
 [ 0.99  1.47  0.    0.5 ]]
```




NMF features

- NMF feature values are non-negative
- Can be used to reconstruct the samples
- ... combine feature values with components

```
In [6]: print(nmf_features)
[[ 0.    0.2 ]
 [ 0.19  0.   ]
 ...
 [ 0.15  0.12]]
```



Reconstruction of a sample

```
In [7]: print(samples[i,:])  
[ 0.12  0.18  0.32  0.14]
```

```
In [8]: print(nmf_features[i,:])  
[ 0.15  0.12]
```

```
0.15 *  
+ 0.12 *
```

```
[[ 0.01  0.   2.13  0.54 ]  
 [ 0.99  1.47  0.   0.5  ]]
```

`model.components_`

reconstruction of sample

```
[ 0.1203  0.1764  0.3195  0.141 ]
```



Sample reconstruction

- Multiply components by feature values, and add up
- Can also be expressed as a product of matrices
- This is the "**Matrix Factorization**" in "NMF"



NMF fits to non-negative data, only

- Word frequencies in each document
- Images encoded as arrays
- Audio spectrograms
- Purchase histories on e-commerce sites
- ... and many more!



UNSUPERVISED LEARNING IN PYTHON

Let's practice!



UNSUPERVISED LEARNING IN PYTHON

**NMF learns
interpretable parts**



Example: NMF learns interpretable parts

- Word-frequency array articles (tf-idf)
- 20,000 scientific articles (rows)
- 800 words (columns)





Applying NMF to the articles

```
In [1]: print(articles.shape)
(20000, 800)
```

```
In [2]: from sklearn.decomposition import NMF
```

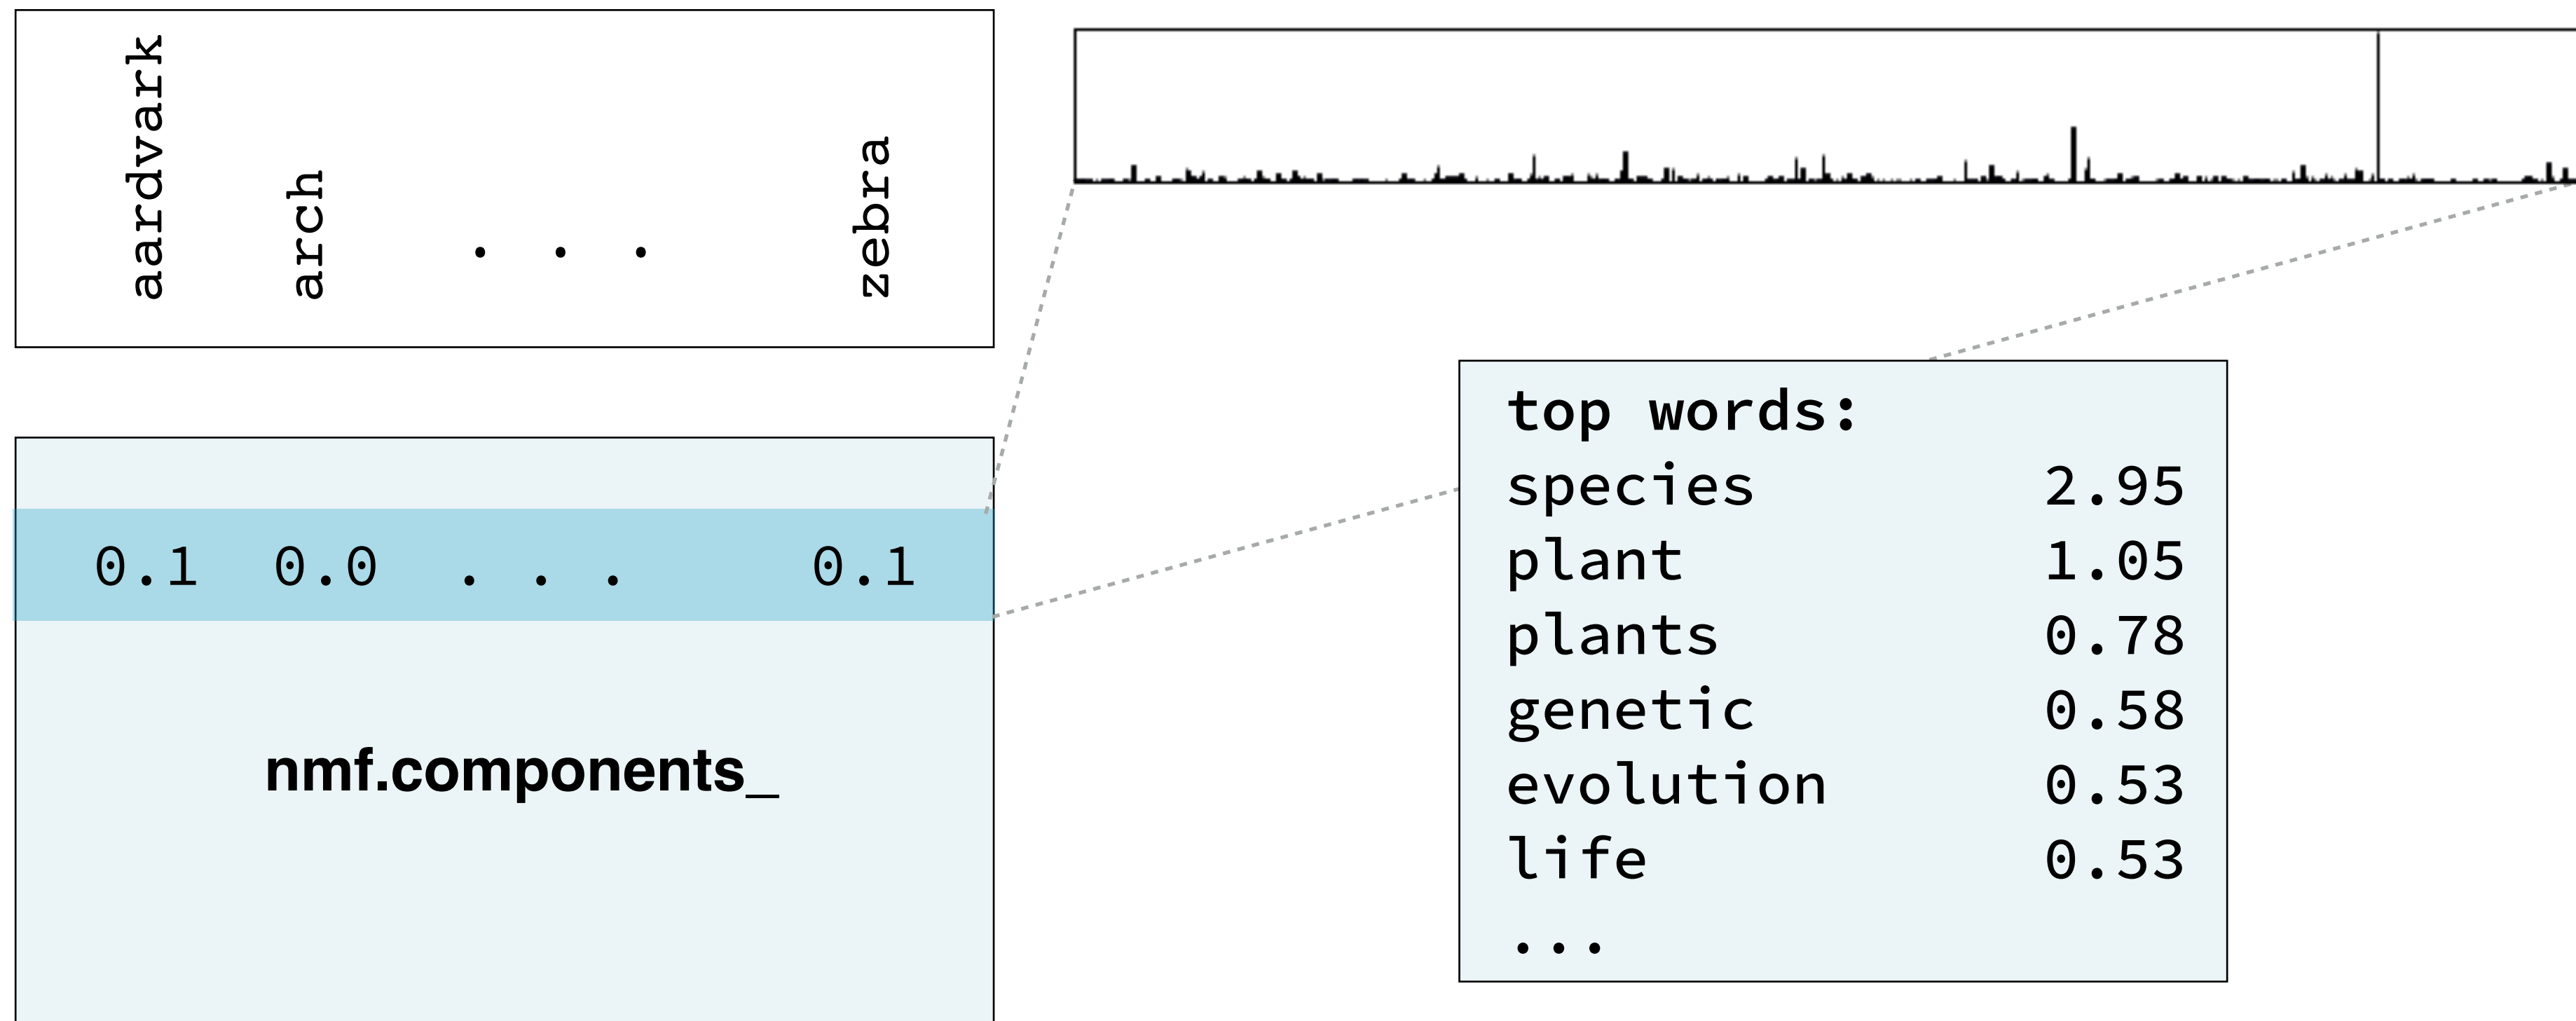
```
In [3]: nmf = NMF(n_components=10)
```

```
In [4]: nmf.fit(articles)
Out[4]: NMF(alpha=0.0, ... )
```

```
In [5]: print(nmf.components_.shape)
(10, 800)
```




NMF components are topics

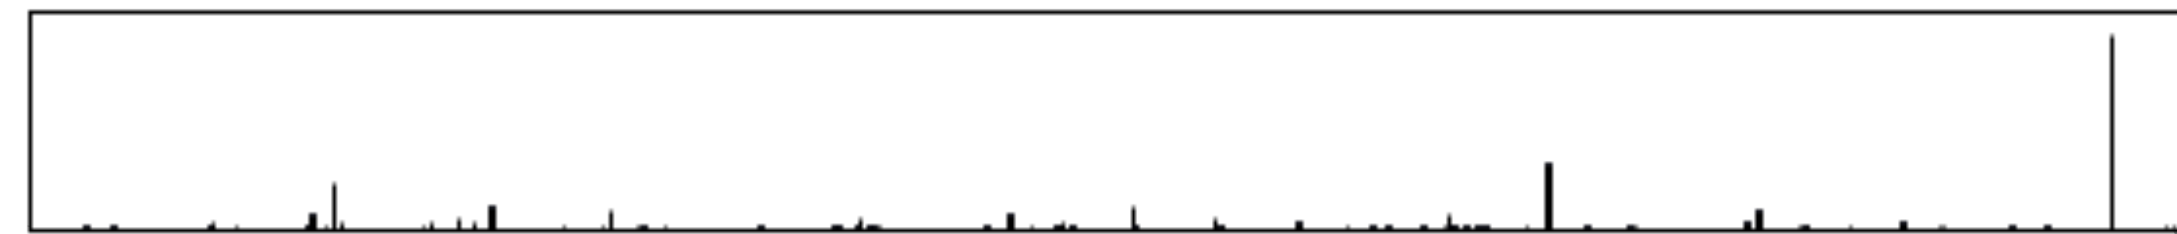




NMF components are topics

aardvark	arch	.	.	.	zebra
----------	------	---	---	---	-------

nmf.components_					
0.0	0.01	.	.	.	0.0



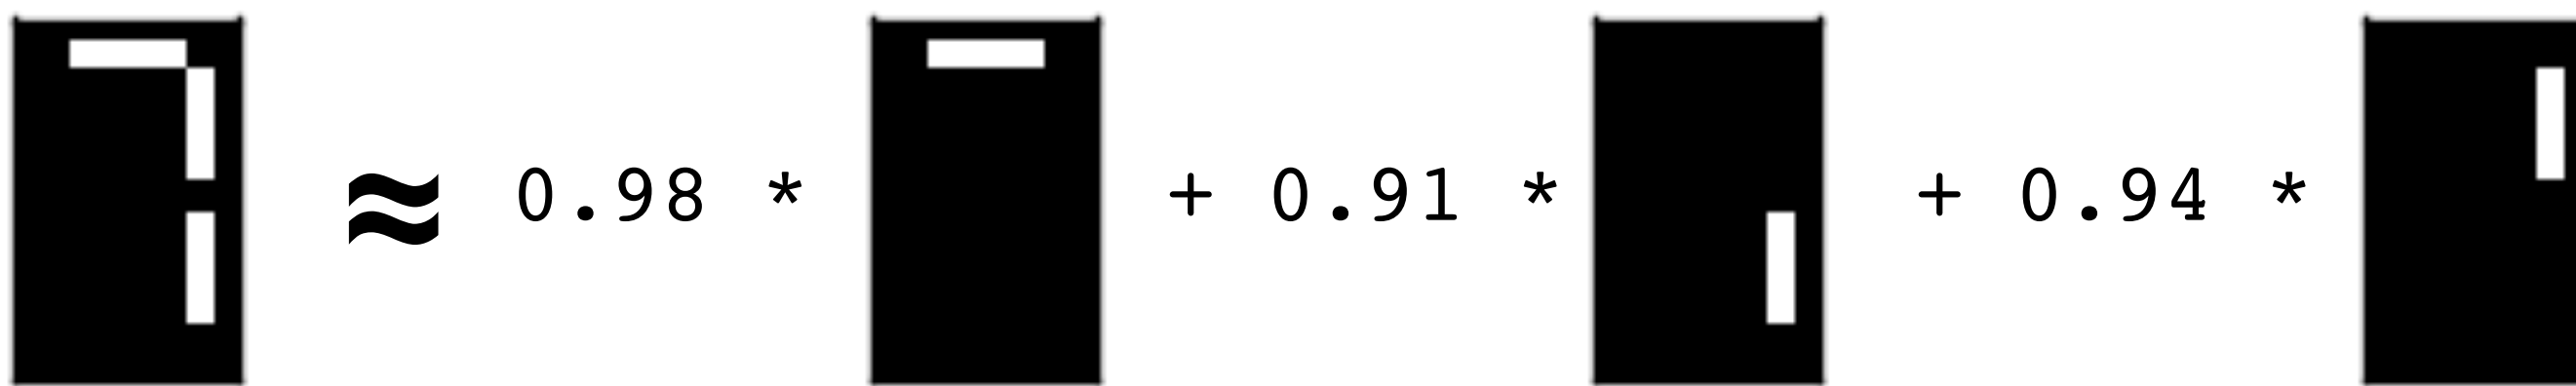
top words:

university	3.57
prof	1.19
college	0.88
department	0.42
education	0.33
school	0.31
...	



NMF components

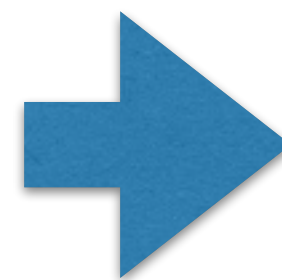
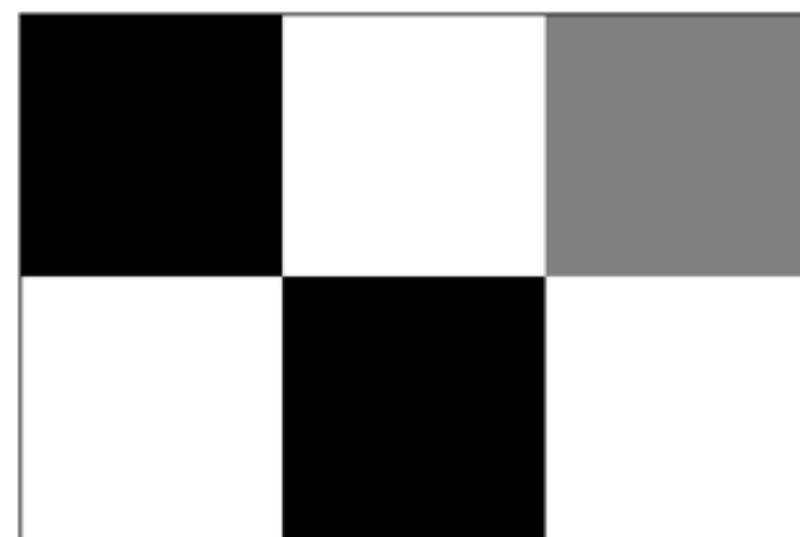
- For documents:
 - NMF components represent topics
 - NMF features combine topics into documents
- For images, NMF components are parts of images


$$\text{Image} \approx 0.98 * \text{Component 1} + 0.91 * \text{Component 2} + 0.94 * \text{Component 3}$$



Grayscale images

- "Grayscale" image = no colors, only shades of gray
- Measure pixel brightness
- Represent with value between 0 and 1 (0 is black)
- Convert to 2D array

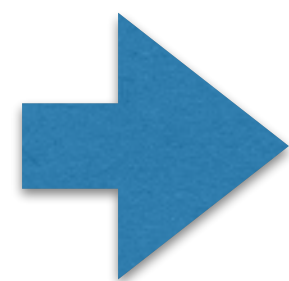


```
[[ 0.  1.  0.5]
 [ 1.  0.  1. ]]
```



Grayscale image example

- An 8x8 grayscale image of the moon, written as an array



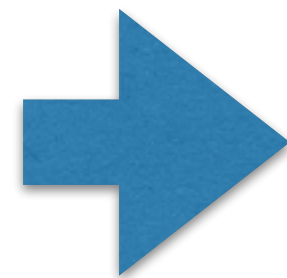
```
[[ 0.  0.  0.  0.  0.  0.  0.  0. ]  
 [ 0.  0.  0.  0.7 0.8 0.  0.  0. ]  
 [ 0.  0.  0.8 0.8 0.9 1.  0.  0. ]  
 [ 0.  0.7 0.9 0.9 1.  1.  1.  0. ]  
 [ 0.  0.8 0.9 1.  1.  1.  1.  0. ]  
 [ 0.  0.  0.9 1.  1.  1.  0.  0. ]  
 [ 0.  0.  0.  0.9 1.  0.  0.  0. ]  
 [ 0.  0.  0.  0.  0.  0.  0.  0. ]]
```



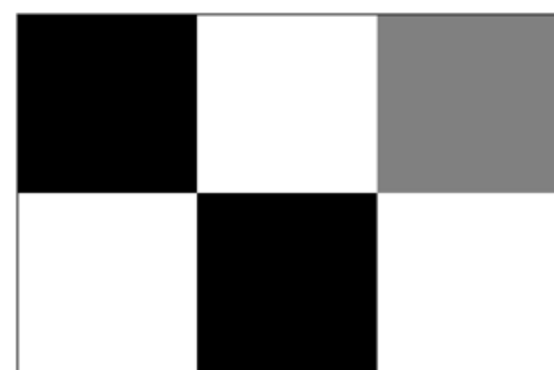
Grayscale images as flat arrays

- Enumerate the entries
- Row-by-row
- From left to right

```
[[ 0.  1.  0.5]  
 [ 1.  0.  1. ]]
```



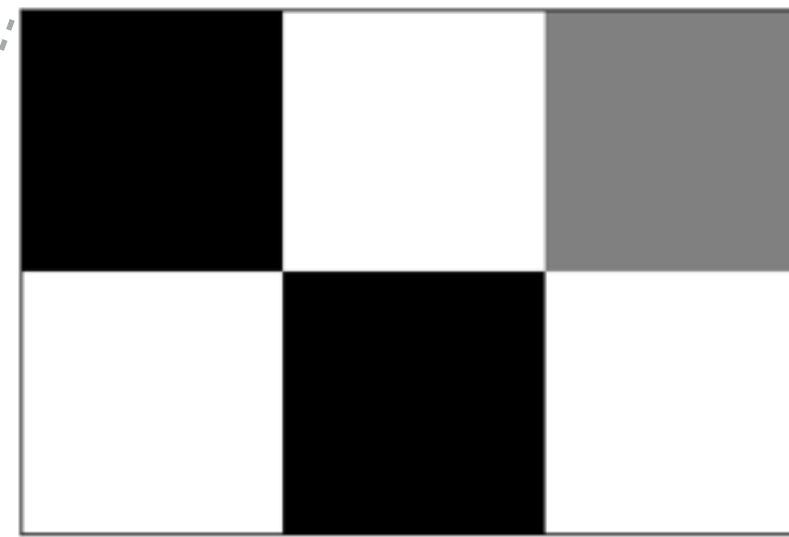
```
[ 0.  1.  0.5  1.  0.  1. ]
```





Encoding a collection of images

- Collection of images of the same size
- Encode as 2D array
- Each row corresponds to an image
- Each column corresponds to a pixel
- ... can apply NMF!



0.	1.	0.5	1.	0.	1.
----	----	-----	----	----	----



Visualizing samples

```
In [1]: print(sample)
[ 0.   1.   0.5  1.   0.   1. ]

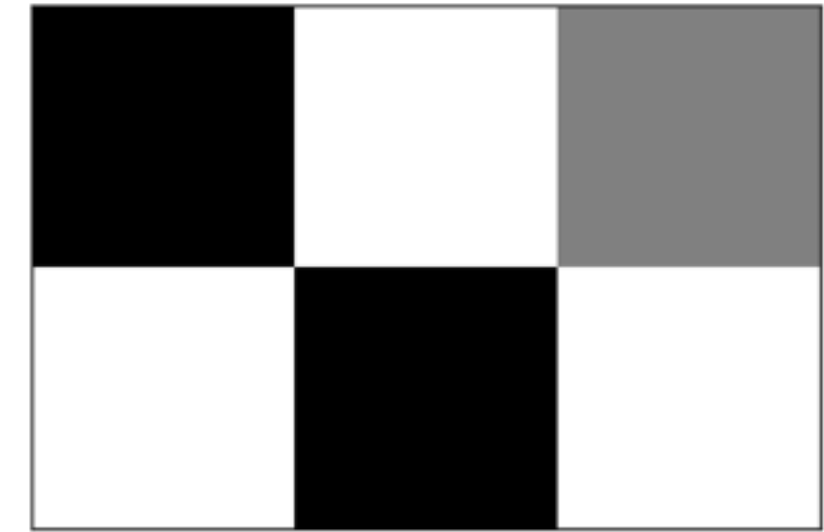
In [2]: bitmap = sample.reshape((2, 3))

In [3]: print(bitmap)
[[ 0.   1.   0.5]
 [ 1.   0.   1. ]]

In [4]: from matplotlib import pyplot as plt

In [5]: plt.imshow(bitmap, cmap='gray', interpolation='nearest')

In [6]: plt.show()
```





UNSUPERVISED LEARNING IN PYTHON

Let's practice!



UNSUPERVISED LEARNING IN PYTHON

Building recommender systems using NMF



Finding similar articles

- Engineer at a large online newspaper
- Task: recommend articles similar to article being read by customer
- Similar articles should have similar topics



Strategy

- Apply NMF to the word-frequency array
- NMF feature values describe the topics
- ... so similar documents have similar NMF feature values
- Compare NMF feature values?



Apply NMF to the word-frequency array

- `articles` is a word frequency array

```
In [1]: from sklearn.decomposition import NMF
```

```
In [2]: nmf = NMF(n_components=6)
```

```
In [3]: nmf_features = nmf.fit_transform(articles)
```

Strategy

- Apply NMF to the word-frequency array
- NMF feature values describe the topics
- ... so similar documents have similar NMF feature values
- Compare NMF feature values?



Versions of articles

- Different versions of the same document have same topic *proportions*
- ... exact feature values may be different!
- E.g. because one version uses many meaningless words
- But all versions lie on the same *line* through the origin

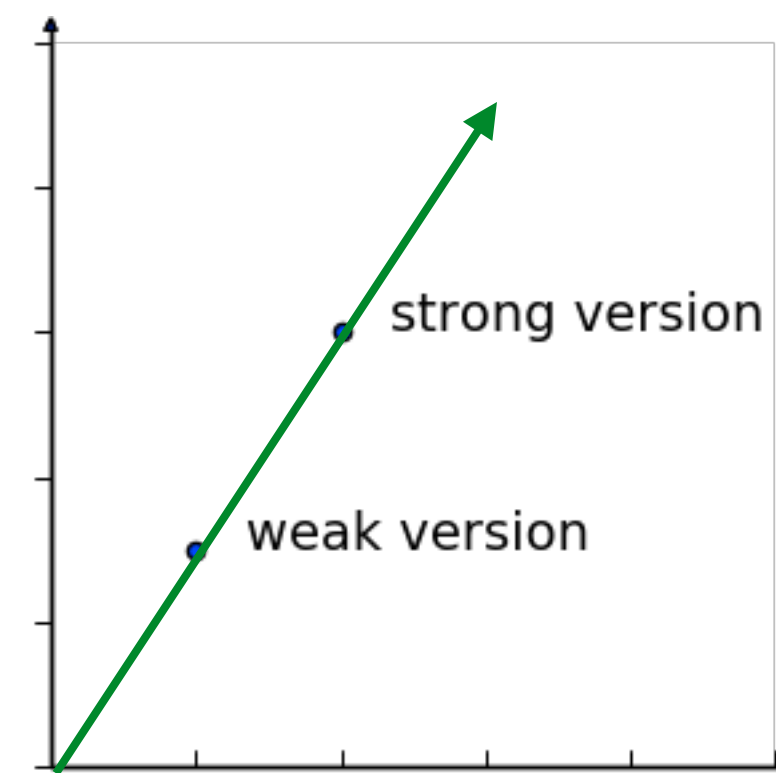
strong version

Dog bites man!
Attack by terrible
canine leaves man
paralyzed...

weak version

You may have heard,
unfortunately it seems
that a dog has perhaps
bitten a man ...

topic:
danger

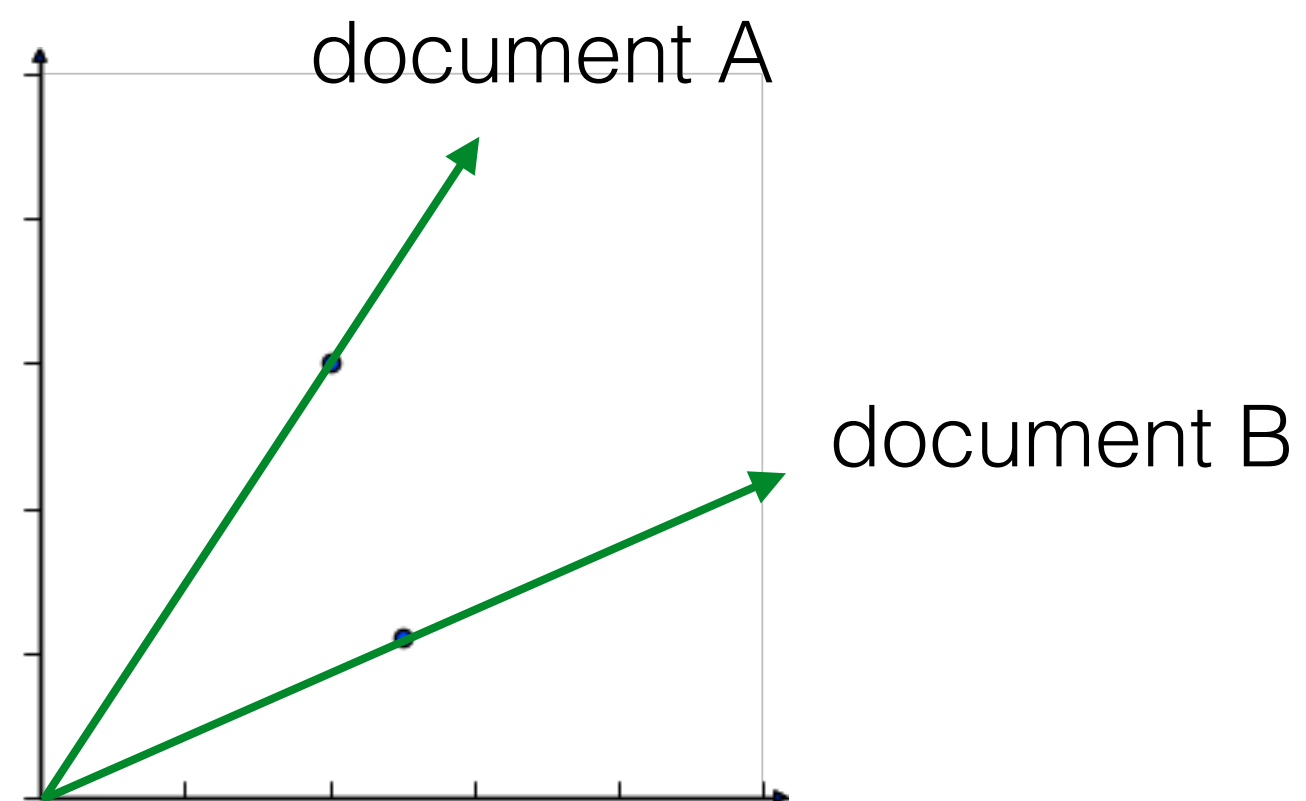


topic:
pets



Cosine similarity

- Uses the angle between the lines
- Higher values means more similar
- Maximum value is 1, when angle is 0°





Calculating the cosine similarities

```
In [4]: from sklearn.preprocessing import normalize

In [5]: norm_features = normalize(nmf_features)

In [6]: current_article = norm_features[23,:] # if has index 23

In [7]: similarities = norm_features.dot(current_article)

In [8]: print(similarities)
[ 0.7150569  0.26349967 0.40210445 ...,  0.70462768  0.20323616
 0.05047817]
```



DataFrames and labels

- Label similarities with the article titles, using a DataFrame
- Titles given as a list: **titles**

```
In [9]: import pandas as pd
```

```
In [10]: norm_features = normalize(nmf_features)
```

```
In [11]: df = pd.DataFrame(norm_features, index=titles)
```

```
In [12]: current_article = df.loc['Dog bites man']
```

```
In [13]: similarities = df.dot(current_article)
```



DataFrames and labels

```
In [14]: print(similarities.nlargest())  
Dog bites man                                1.000000  
Hound mauls cat                             0.979946  
Pets go wild!                               0.979708  
Dachshunds are dangerous                    0.949641  
Our streets are no longer safe               0.900474  
dtype: float64
```



UNSUPERVISED LEARNING IN PYTHON

Let's practice!



UNSUPERVISED LEARNING IN PYTHON

Building recommender systems using NMF



Finding similar articles

- Engineer at a large online newspaper
- Task: recommend articles similar to article being read by customer
- Similar articles should have similar topics

Strategy

- Apply NMF to the word-frequency array
- NMF feature values describe the topics
- ... so similar documents have similar NMF feature values
- Compare NMF feature values?



Apply NMF to the word-frequency array

- `articles` is a word frequency array

```
In [1]: from sklearn.decomposition import NMF
```

```
In [2]: nmf = NMF(n_components=6)
```

```
In [3]: nmf_features = nmf.fit_transform(articles)
```




Strategy

- Apply NMF to the word-frequency array
- NMF feature values describe the topics
- ... so similar documents have similar NMF feature values
- Compare NMF feature values?



Versions of articles

- Different versions of the same document have same topic *proportions*
- ... exact feature values may be different!
- E.g. because one version uses many meaningless words
- But all versions lie on the same *line* through the origin

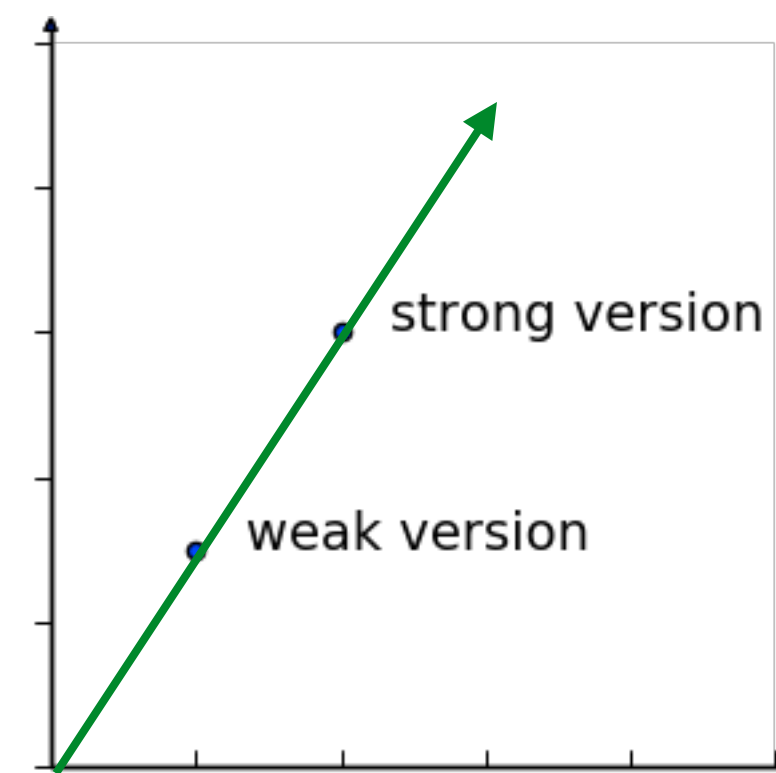
strong version

Dog bites man!
Attack by terrible
canine leaves man
paralyzed...

weak version

You may have heard,
unfortunately it seems
that a dog has perhaps
bitten a man ...

topic:
danger

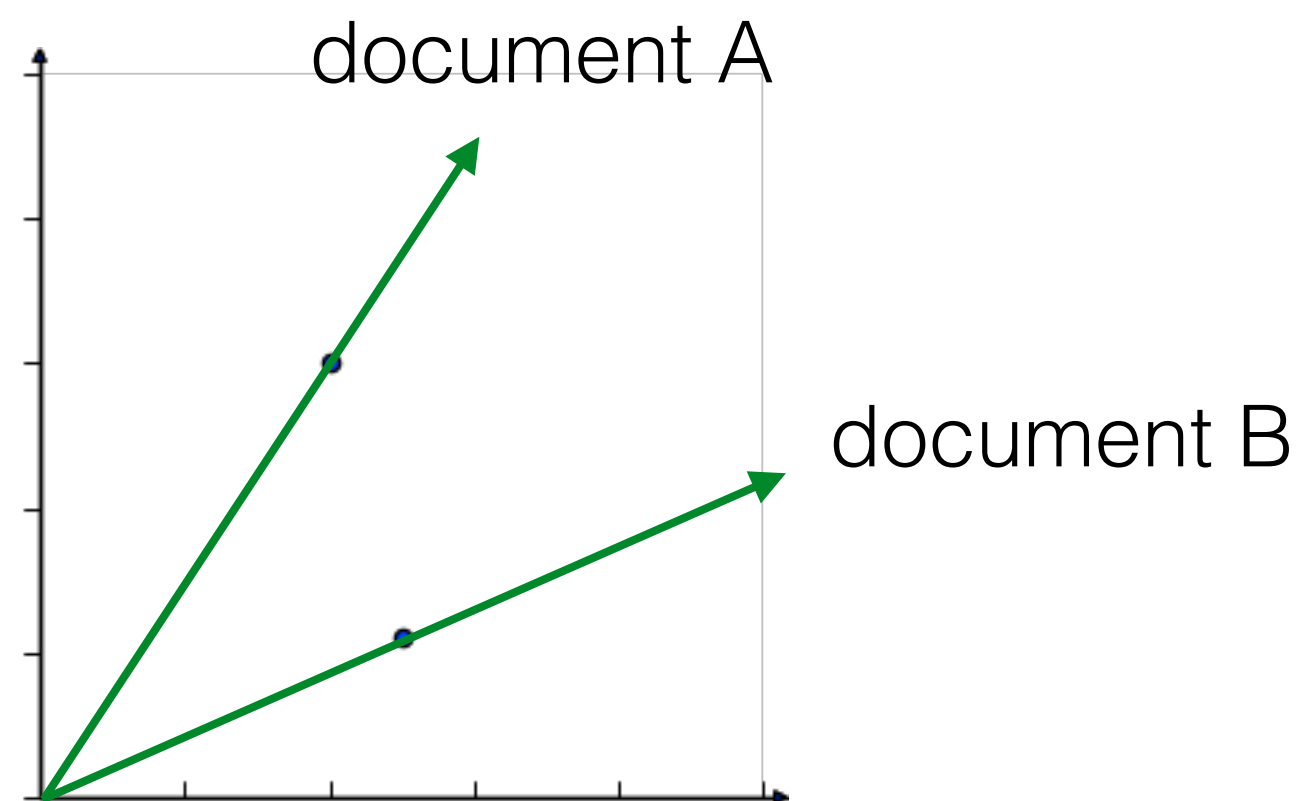


topic:
pets



Cosine similarity

- Uses the angle between the lines
- Higher values means more similar
- Maximum value is 1, when angle is 0°





Calculating the cosine similarities

```
In [4]: from sklearn.preprocessing import normalize

In [5]: norm_features = normalize(nmf_features)

In [6]: current_article = norm_features[23,:] # if has index 23

In [7]: similarities = norm_features.dot(current_article)

In [8]: print(similarities)
[ 0.7150569  0.26349967 0.40210445 ...,  0.70462768  0.20323616
 0.05047817]
```



DataFrames and labels

- Label similarities with the article titles, using a DataFrame
- Titles given as a list: **titles**

```
In [9]: import pandas as pd
```

```
In [10]: norm_features = normalize(nmf_features)
```

```
In [11]: df = pd.DataFrame(norm_features, index=titles)
```

```
In [12]: current_article = df.loc['Dog bites man']
```

```
In [13]: similarities = df.dot(current_article)
```



DataFrames and labels

```
In [14]: print(similarities.nlargest())  
Dog bites man                                1.000000  
Hound mauls cat                             0.979946  
Pets go wild!                               0.979708  
Dachshunds are dangerous                    0.949641  
Our streets are no longer safe              0.900474  
dtype: float64
```



UNSUPERVISED LEARNING IN PYTHON

Let's practice!